

Nour2.foodboost

January 25, 2023

1 import libraries

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import random
from sklearn.metrics import classification_report, confusion_matrix, roc_curve,
    mean_squared_error, accuracy_score, recall_score, precision_score, f1_score
from sklearn.metrics import roc_auc_score, mean_absolute_error, make_scorer
from sklearn.model_selection import cross_val_score, train_test_split, KFold,
    RandomizedSearchCV, GridSearchCV
from sklearn.linear_model import Lasso
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import make_column_transformer
from sklearn.ensemble import GradientBoostingClassifier
from pipetorch import DFrame
from sklearn.svm import SVC, LinearSVC
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from math import sqrt
from scipy import stats
from seaborn import load_dataset, pairplot
from sklearn import tree
from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier
from pipetorch.evaluate.study import Study, optuna
```

2 uiteindelijk simulatie inladen

```
[10]: sim_groot= pd.read_csv('Uiteindelijk dataset food Heel Groot.csv')
sim= sim_groot.sample(frac=1).reset_index(drop=True)
sim
```

```

[10]:      favo_Aziatische noedelsalade met tofu \
0          0.0
1          0.0
2          1.0
3          0.0
4          0.0
...          ...
69995      0.0
69996      0.0
69997      0.0
69998      1.0
69999      0.0

      favo_Aziatische roerbak met varkenssaté \
0          0.0
1          0.0
2          0.0
3          0.0
4          0.0
...          ...
69995      0.0
69996      1.0
69997      0.0
69998      0.0
69999      0.0

      favo_Champignonragout met groene asperges \
0          0.0
1          0.0
2          0.0
3          0.0
4          0.0
...          ...
69995      0.0
69996      0.0
69997      0.0
69998      0.0
69999      0.0

      favo_Chinese balletjes in zoetzure saus \
0          0.0
1          0.0
2          0.0
3          0.0
4          0.0
...          ...
69995      0.0

```

69996	0.0
69997	0.0
69998	0.0
69999	0.0

	favo_Eenpansnoedels met vega kip en wokgroente \
0	0.0
1	0.0
2	0.0
3	0.0
4	1.0
...	...
69995	1.0
69996	0.0
69997	0.0
69998	0.0
69999	1.0

	favo_Falafel-quinoaburger met zuivelspread, slamelange, rode biet en komkommer \
0	0.0
1	1.0
2	0.0
3	0.0
4	1.0
...	...
69995	0.0
69996	0.0
69997	0.0
69998	0.0
69999	0.0

	favo_Gebakken gnocchi met spekjes en spruitjes \
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0
...	...
69995	0.0
69996	0.0
69997	0.0
69998	0.0
69999	0.0

	favo_Groenten in Thaise currysaus \
0	1.0

1	0.0
2	0.0
3	0.0
4	0.0
...	...
69995	0.0
69996	0.0
69997	0.0
69998	0.0
69999	0.0

	favo_Hoisinnoedels met gehakt en groene groenten \
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0
...	...
69995	0.0
69996	0.0
69997	0.0
69998	0.0
69999	0.0

	favo_Kip roerbak met sperziebonen en rijst ... \
0	0.0 ...
1	0.0 ...
2	0.0 ...
3	0.0 ...
4	0.0 ...
...
69995	0.0 ...
69996	0.0 ...
69997	1.0 ...
69998	0.0 ...
69999	0.0 ...

	voorspel_Spinazie-couscoussalade met ei	voorspel_Spinaziecurry met ei \
0	0.0	0.0
1	0.0	0.0
2	1.0	0.0
3	0.0	0.0
4	0.0	0.0
...
69995	0.0	0.0
69996	0.0	0.0
69997	0.0	0.0

69998	0.0	0.0
69999	0.0	0.0

	voorspel_Supersnel rijstpannetje met tuinbonen \
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0
...	...
69995	0.0
69996	0.0
69997	0.0
69998	0.0
69999	0.0

	voorspel_Tomaten-tuinbonensalade \
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0
...	...
69995	0.0
69996	0.0
69997	0.0
69998	0.0
69999	0.0

	voorspel_Vegan curry met pitabroodjes \
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0
...	...
69995	0.0
69996	0.0
69997	0.0
69998	0.0
69999	0.0

	voorspel_Vegetarische wokschotel met bloemkoolrijst en sweet chilisaus \
0	0.0
1	0.0
2	0.0

3	0.0
4	0.0
...	...
69995	0.0
69996	0.0
69997	0.0
69998	0.0
69999	0.0

voorspel_Volkoren pasta met kip, Chinese kool, gember en tahin \	
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0
...	...
69995	0.0
69996	0.0
69997	0.0
69998	0.0
69999	0.0

voorspel_Wokschotel met verse-kaasblokjes \	
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0
...	...
69995	0.0
69996	0.0
69997	0.0
69998	0.0
69999	0.0

voorspel_Zomerse chili con carne target		
0	0.0	0
1	0.0	1
2	0.0	1
3	0.0	1
4	0.0	1
...
69995	0.0	1
69996	0.0	0
69997	0.0	1
69998	0.0	0
69999	0.0	0

[70000 rows x 91 columns]

```
[11]: X = sim.drop(sim.columns[-1], axis = 1).values
      y = sim[sim.columns[-1]].values
      X_train, X_test, y_train, y_test = \
          train_test_split(X, y, test_size=0.2)
```

Dit is een classificatie probleem. Daardoor zijn de volgende machine learning modellen toegepast: Logistic Regression, KNeighbors Classifier, RandomForest Classifier, Decision Tree en Support Vector Classifier.

3 Logistic Regression

parameter `class_weight` gebruikt om de data te balanceren

```
[12]: logreg = LogisticRegression(class_weight='balanced')
      logreg.fit(X_train, y_train)
      pred_y= logreg.predict(X_test)
      pred_y_prob = logreg.predict_proba(X_test)[:,:1]
      print(confusion_matrix(y_test, pred_y))
      print(classification_report(y_test, pred_y))
      precision = precision_score(y_test, pred_y)
      recall = recall_score(y_test, pred_y)
      accuracy = accuracy_score(y_test, pred_y)
      fpr_logreg1, tpr_logreg1, thresholds = roc_curve(y_test, pred_y_prob)
      auc_score_logreg1 = roc_auc_score(y_test, pred_y_prob)

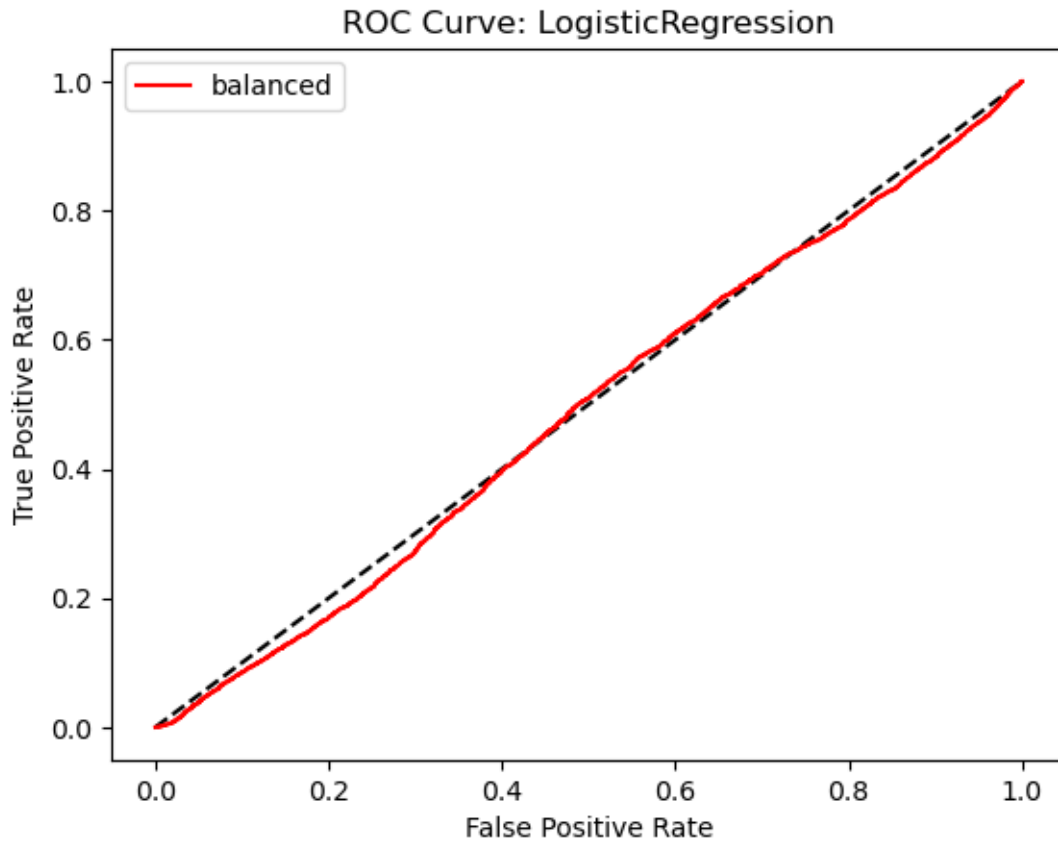
      print(f"accuracy: {accuracy}" )
      print(f"precision: {precision}" )
      print(f"recall: {recall}")
      plt.plot([0, 1], [0, 1], 'k--')
      plt.plot(fpr_logreg1, tpr_logreg1, color = 'red', label = 'balanced')
      plt.xlabel('False Positive Rate')
      plt.ylabel('True Positive Rate')
      plt.title('ROC Curve: LogisticRegression')
      plt.legend()
      plt.show()
```

```
[[3434 3561]
 [3364 3641]]
```

	precision	recall	f1-score	support
0	0.51	0.49	0.50	6995
1	0.51	0.52	0.51	7005

accuracy			0.51	14000
macro avg	0.51	0.51	0.51	14000
weighted avg	0.51	0.51	0.51	14000

accuracy: 0.5053571428571428
precision: 0.5055540127742294
recall: 0.5197715917201998



4 KNeighbors Classifier

```
[13]: knn = KNeighborsClassifier(n_neighbors = 50)
      knn.fit(X_train, y_train)
      print(knn.score(X_test, y_test))
      pred_y = knn.predict(X_test)
      pred_y_prob = knn.predict_proba(X_test)[:,:1]
      print(confusion_matrix(y_test, pred_y))
      print(classification_report(y_test, pred_y))
      precision = precision_score(y_test, pred_y)
```



```

recall = recall_score(y_test, pred_y)
accuracy = accuracy_score(y_test, pred_y)
fpr_knn1, tpr_knn1, thresholds = roc_curve(y_test, pred_y_prob)
auc_score_knn1 = roc_auc_score(y_test, pred_y_prob)

print(f"accuracy: {accuracy}" )
print(f"precision: {precision}" )
print(f"recall: {recall}")
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr_knn1, tpr_knn1, color = 'red', label='n_neighbors = 50')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve: KNeighbors')
plt.legend()
plt.show()

```

0.7999285714285714

[[5415 1580]

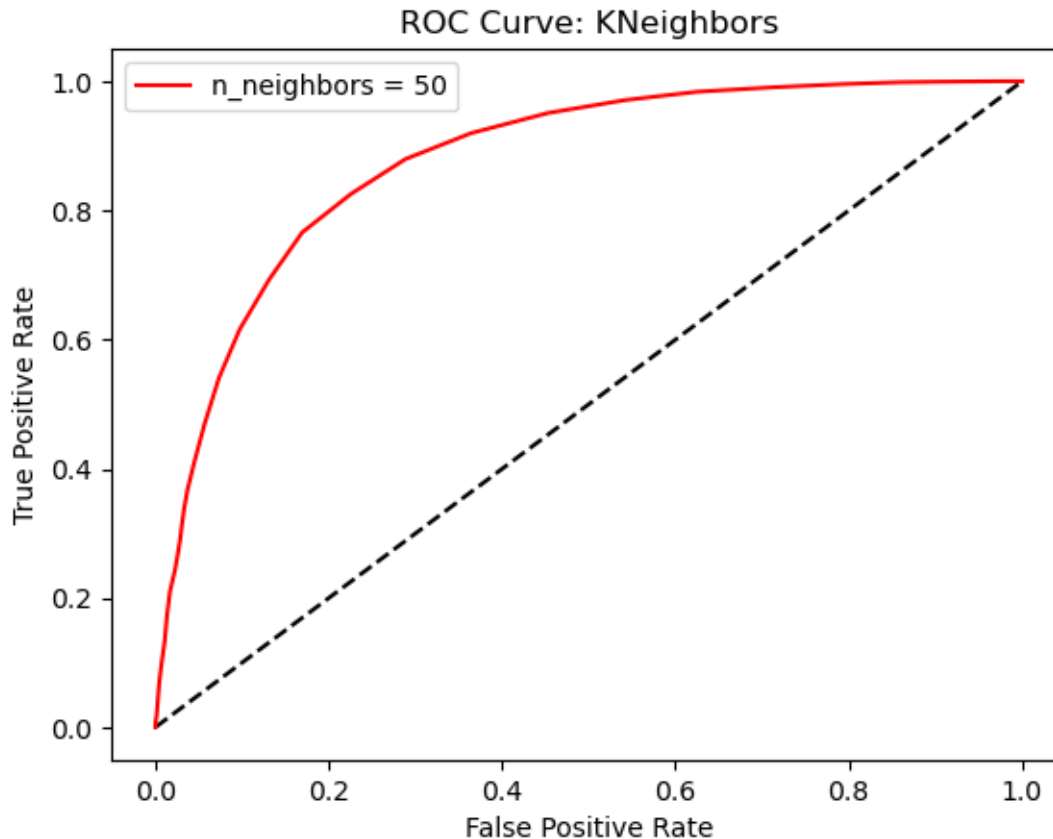
[1221 5784]]

	precision	recall	f1-score	support
0	0.82	0.77	0.79	6995
1	0.79	0.83	0.81	7005
accuracy			0.80	14000
macro avg	0.80	0.80	0.80	14000
weighted avg	0.80	0.80	0.80	14000

accuracy: 0.7999285714285714

precision: 0.7854426941879413

recall: 0.8256959314775161



5 RandomForestClassifier

```
[14]: rfc= RandomForestClassifier(max_depth=22, n_estimators=180,
                                random_state=42)
rfc.fit(X_train, y_train)
pred_y = rfc.predict(X_test)
pred_y_prob = rfc.predict_proba(X_test)[:,:1]
print(confusion_matrix(y_test, pred_y))
print(classification_report(y_test, pred_y))
precision = precision_score(y_test, pred_y)
recall = recall_score(y_test, pred_y)
accuracy = accuracy_score(y_test, pred_y)
fpr_rfc1, tpr_rfc1, thresholds = roc_curve(y_test, pred_y_prob)
auc_score_rfc1 = roc_auc_score(y_test, pred_y_prob)

print(f"accuracy: {accuracy}" )
print(f"precision: {precision}" )
print(f"recall: {recall}")
```


```
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr_rfc1, tpr_rfc1, color = 'red')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve: RandomForestClassifier')
plt.legend()
plt.show()
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

```
[[5763 1232]
 [ 434 6571]]
```

	precision	recall	f1-score	support
0	0.93	0.82	0.87	6995
1	0.84	0.94	0.89	7005
accuracy			0.88	14000
macro avg	0.89	0.88	0.88	14000
weighted avg	0.89	0.88	0.88	14000

```
accuracy: 0.881
precision: 0.8421120082019736
recall: 0.9380442541042113
```



6 DecisionTree

```
[15]: clf = DecisionTreeClassifier(class_weight='balanced', splitter='best')
      clf.fit(X_train, y_train)
      pred_y = clf.predict(X_test)
      pred_y_prob = clf.predict_proba(X_test)[:,:1]
      print(confusion_matrix(y_test, pred_y))
      print(classification_report(y_test, pred_y))
      precision = precision_score(y_test, pred_y)
      recall = recall_score(y_test, pred_y)
      accuracy = accuracy_score(y_test, pred_y)
      fpr_clf1, tpr_clf1, thresholds = roc_curve(y_test, pred_y_prob)
      auc_score_clf1 = roc_auc_score(y_test, pred_y_prob)


      print(f"accuracy: {accuracy}")
      print(f"precision: {precision}")
      print(f"recall: {recall}")
      plt.plot([0, 1], [0, 1], 'k--')
```

```
plt.plot(fpr_clf1, tpr_clf1, color = 'red', label='class_weight=balanced, splitter=best')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve: DecisionTreeClassifier')
plt.legend()
plt.show()
```

```
[[5947 1048]
 [ 892 6113]]
```

	precision	recall	f1-score	support
0	0.87	0.85	0.86	6995
1	0.85	0.87	0.86	7005
accuracy			0.86	14000
macro avg	0.86	0.86	0.86	14000
weighted avg	0.86	0.86	0.86	14000

```
accuracy: 0.8614285714285714
precision: 0.8536517246194666
recall: 0.8726623840114204
```




7 Support Vector Classifier

```
[16]: svc= SVC(C=10, gamma=0.1)
svc.fit(X_train, y_train)
pred_y = svc.predict(X_test)
pred_y_prob = svc.decision_function(X_test)
print(confusion_matrix(y_test, pred_y))
precision = precision_score(y_test, pred_y)
recall = recall_score(y_test, pred_y)
accuracy = accuracy_score(y_test, pred_y)
fpr_svc1, tpr_svc1, thresholds = roc_curve(y_test, pred_y_prob)
auc_score_svc1 = roc_auc_score(y_test, pred_y_prob)

print(f"accuracy: {accuracy}" )
print(f"precision: {precision}" )
print(f"recall: {recall}")
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr_svc1, tpr_svc1, color = 'red')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve: SVC')
plt.legend()
plt.show()
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

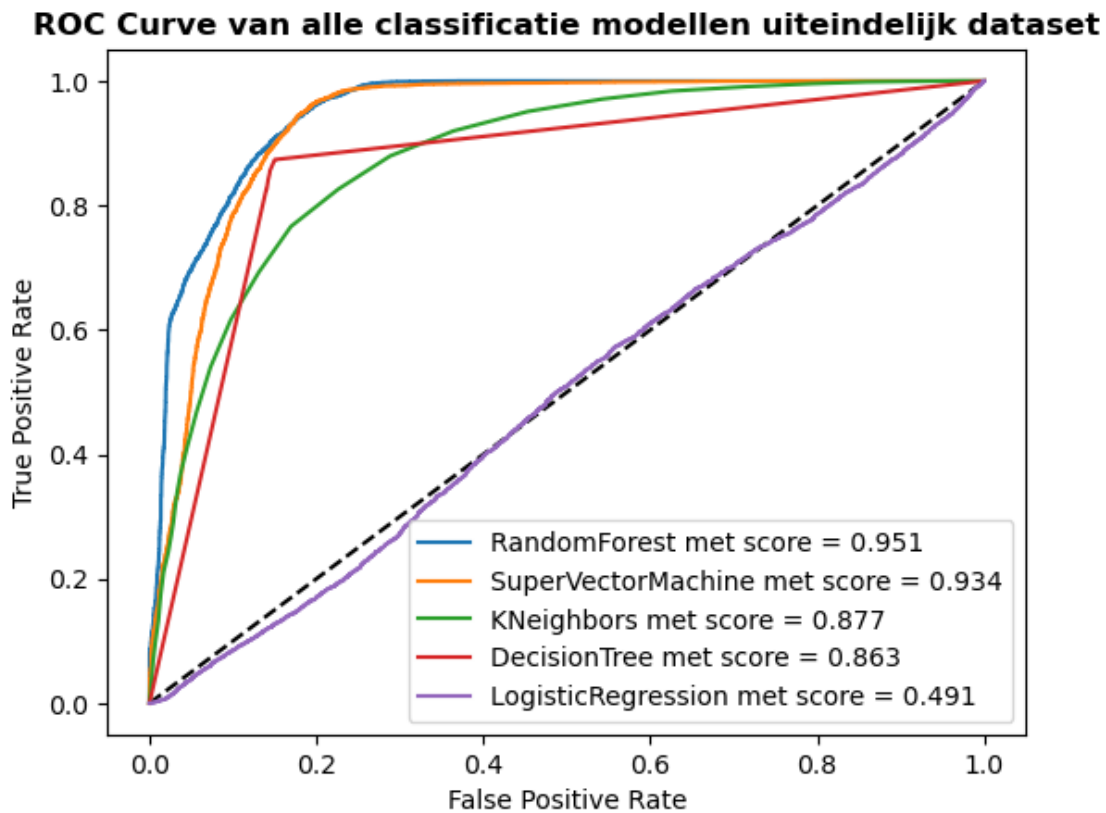
```
[[5785 1210]
 [ 459 6546]]
accuracy: 0.8807857142857143
precision: 0.8439917483238782
recall: 0.934475374732334
```



8 ROC Curve van alle modellen

```
[17]: plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr_rfc1, tpr_rfc1, label = 'RandomForest met score = %0.3f' % auc_score_rfc1)
plt.plot(fpr_svc1, tpr_svc1, label = 'SuperVectorMachine met score = %0.3f' % auc_score_svc1)
plt.plot(fpr_knn1, tpr_knn1, label = 'KNeighbors met score = %0.3f' % auc_score_knn1)
plt.plot(fpr_clf1, tpr_clf1, label = 'DecisionTree met score = %0.3f' % auc_score_clf1)
plt.plot(fpr_logreg1, tpr_logreg1, label = 'LogisticRegression met score = %0.3f' % auc_score_logreg1)
titel = 'ROC Curve van alle classificatie modellen uiteindelijk dataset'
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title(titel, fontweight="bold")
plt.legend()
```

```
plt.savefig('ROC Curve alle modellen grote dataset.png')
plt.show()
```



9 optuna study om de beste parameters te achterhalen

```
[19]: X_train, X_valid, y_train, y_valid= train_test_split(sim.
      ↪drop(columns='target'), sim.target, test_size=0.2)
study= optuna.create_study()
def trial(trial):
    depth= trial.suggest_categorical('max_depth', range(1,51))
    n= trial.suggest_categorical('n_estimators', range(1,1001))
    model= RandomForestClassifier(max_depth= depth, n_estimators= n)
    model.fit(X_train, y_train)
    pre_valid= precision_score(y_valid, model.predict(X_valid))
    return pre_valid
study.optimize(trial, n_trials=20)
```

[I 2023-01-21 16:47:11,276] A new study created in memory with name:
no-name-077dd544-71b1-4f30-858a-2c0ff1014dcc

[I 2023-01-21 16:47:54,614] Trial 0 finished with value:
0.8461035564853556 and parameters: {'max_depth': 48, 'n_estimators': 631}. Best
is trial 0 with value: 0.8461035564853556.

[I 2023-01-21 16:48:22,975] Trial 1 finished with value:
0.8416376979528776 and parameters: {'max_depth': 15, 'n_estimators': 912}. Best
is trial 1 with value: 0.8416376979528776.

[I 2023-01-21 16:48:53,533] Trial 2 finished with value:
0.8446614416034948 and parameters: {'max_depth': 20, 'n_estimators': 789}. Best
is trial 1 with value: 0.8416376979528776.

[I 2023-01-21 16:49:13,755] Trial 3 finished with value:
0.838097679784643 and parameters: {'max_depth': 15, 'n_estimators': 662}. Best
is trial 3 with value: 0.838097679784643.

[I 2023-01-21 16:49:45,492] Trial 4 finished with value:
0.8469400957433044 and parameters: {'max_depth': 32, 'n_estimators': 559}. Best
is trial 3 with value: 0.838097679784643.

[I 2023-01-21 16:50:01,194] Trial 5 finished with value:
0.8397486252945797 and parameters: {'max_depth': 10, 'n_estimators': 723}. Best
is trial 3 with value: 0.838097679784643.

[I 2023-01-21 16:51:00,025] Trial 6 finished with value:
0.8466153645493674 and parameters: {'max_depth': 46, 'n_estimators': 916}. Best
is trial 3 with value: 0.838097679784643.

[I 2023-01-21 16:51:22,230] Trial 7 finished with value:
0.8471428571428572 and parameters: {'max_depth': 37, 'n_estimators': 389}. Best
is trial 3 with value: 0.838097679784643.

[I 2023-01-21 16:51:32,615] Trial 8 finished with value:
0.8418748390419779 and parameters: {'max_depth': 16, 'n_estimators': 330}. Best
is trial 3 with value: 0.838097679784643.

[I 2023-01-21 16:51:59,339] Trial 9 finished with value:
0.839677047289504 and parameters: {'max_depth': 16, 'n_estimators': 854}. Best
is trial 3 with value: 0.838097679784643.

[I 2023-01-21 16:52:51,854] Trial 10 finished with value:
0.8458952811893988 and parameters: {'max_depth': 33, 'n_estimators': 975}. Best
is trial 3 with value: 0.838097679784643.

[I 2023-01-21 16:53:06,183] Trial 11 finished with value:
0.8421748908585792 and parameters: {'max_depth': 9, 'n_estimators': 662}. Best
is trial 3 with value: 0.838097679784643.

[I 2023-01-21 16:53:11,349] Trial 12 finished with value:
0.8334820859365039 and parameters: {'max_depth': 15, 'n_estimators': 172}. Best
is trial 12 with value: 0.8334820859365039.

[I 2023-01-21 16:53:14,344] Trial 13 finished with value:
0.844688070999739 and parameters: {'max_depth': 15, 'n_estimators': 101}. Best
is trial 12 with value: 0.8334820859365039.

[I 2023-01-21 16:53:34,283] Trial 14 finished with value:
0.842287472752917 and parameters: {'max_depth': 15, 'n_estimators': 670}. Best
is trial 12 with value: 0.8334820859365039.

[I 2023-01-21 16:54:05,561] Trial 15 finished with value:
0.8432605905006418 and parameters: {'max_depth': 26, 'n_estimators': 679}. Best
is trial 12 with value: 0.8334820859365039.

[I 2023-01-21 16:54:19,059] Trial 16 finished with value: 0.844253770150806 and parameters: {'max_depth': 42, 'n_estimators': 222}. Best is trial 12 with value: 0.8334820859365039.

[I 2023-01-21 16:54:46,194] Trial 17 finished with value: 0.8459358495427025 and parameters: {'max_depth': 24, 'n_estimators': 625}. Best is trial 12 with value: 0.8334820859365039.

[I 2023-01-21 16:54:55,015] Trial 18 finished with value: 0.844269691891195 and parameters: {'max_depth': 31, 'n_estimators': 170}. Best is trial 12 with value: 0.8334820859365039.

[I 2023-01-21 16:55:51,417] Trial 19 finished with value: 0.8469148659203333 and parameters: {'max_depth': 43, 'n_estimators': 911}. Best is trial 12 with value: 0.8334820859365039.

10 verschil tussen de eerste randomforestclassifier en die uit optuna study komt

```
[21]: best= RandomForestClassifier(max_depth= 48, n_estimators= 631)
best.fit(X_train, y_train)
pred_y = best.predict(X_test)
pred_y_prob = best.predict_proba(X_test)[: ,1]
print(confusion_matrix(y_test, pred_y))
print(classification_report(y_test, pred_y))
precision = precision_score(y_test, pred_y)
recall = recall_score(y_test, pred_y)
accuracy = accuracy_score(y_test, pred_y)
fpr_rfc_best, tpr_rfc_best, thresholds_best = roc_curve(y_test, pred_y_prob)
auc_score_rfc_best = roc_auc_score(y_test, pred_y_prob)

print(f"accuracy: {accuracy}" )
print(f"precision: {precision}" )
print(f"recall: {recall}")
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr_rfc1, tpr_rfc1, color = 'red', label = 'RandomForest met score = %0.3f' % auc_score_rfc1)
plt.plot(fpr_rfc_best, tpr_rfc_best, color = 'blue', label = 'RandomForest met score = %0.3f' % auc_score_rfc_best)

plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve: RandomForestClassifier best')
plt.legend()
plt.show()
```


/opt/jupyterhub/anaconda/lib/python3.9/site-packages/sklearn/base.py:450:
UserWarning: X does not have valid feature names, but RandomForestClassifier was

```
fitted with feature names
warnings.warn(
/opt/jupyterhub/anaconda/lib/python3.9/site-packages/sklearn/base.py:450:
UserWarning: X does not have valid feature names, but RandomForestClassifier was
fitted with feature names
warnings.warn(
```

```
[[6590  405]
 [ 133 6872]]
```

	precision	recall	f1-score	support
0	0.98	0.94	0.96	6995
1	0.94	0.98	0.96	7005
accuracy			0.96	14000
macro avg	0.96	0.96	0.96	14000
weighted avg	0.96	0.96	0.96	14000

```
accuracy: 0.9615714285714285
precision: 0.944345197196647
recall: 0.9810135617416131
```



uit al die modellen en figuren blijkt RandomForestClassifier van de optuna study het beste te zijn.

[]: